

Detección de Hate Speech en Twitter





MOTIVACIÓN

El discurso del odio en redes sociales conduce a la normalización de la discriminación, la intolerancia y a actitudes y comportamientos hostiles.





MOTIVACIÓN



A Genocide Incited on Facebook, With Posts From Myanmar's Military

A/HRC/39/64

Advance Edited Version

Distr.: General
12 September 2018

Original: English

Human Rights Council

Thirty-ninth session
10–28 September 2018

Agenda item 4

Human rights situations that require the Council's attention

**Report of the independent international fact-finding mission
on Myanmar***

“Facebook has been a useful
instrument for those seeking
to spread hate ...”

Hate speech on Twitter predicts frequency of real-life hate crimes

NYU researchers turn to artificial intelligence to show the links between
online hate and offline violence in 100 cities



Sri Lanka accuses Facebook over hate speech after deadly riots


**The
Guardian**




Dataset



Dataset desbalanceado



Dataset balanceado




Hate speech: Cualquier comunicación que desacredite a una persona o un grupo en función de alguna característica, como raza, color, origen étnico, género, orientación sexual, nacionalidad, religión u otra característica.

Nockleby (2002)

count	hate_speech	offensive_language	neither	clase	tweet	
0	3	0	0	3	2	!!! RT @mayasolovely: As a woman you shouldn't complain about cleaning up your house. & as a man you should always take the trash out...
1	3	0	3	0	1	!!!! RT @mleew17: boy dats cold...tyga dwn bad for cuffin dat hoe in the 1st place!!
2	3	0	3	0	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby4life: You ever fuck a bitch and she start to cry? You be confused as shit
3	3	0	2	1	1	!!!!!!! RT @C_G_Anderson: @viva_based she look like a tranny
4	6	0	6	0	1	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you hear about me might be true or it might be faker than the bitch who told it to ya 



Preprocesamiento



N-gram and tokenization

```
count_vectorizer.fit_transform(["she look like a tranny"])  
print(count_vectorizer.get_feature_names())
```

```
['like', 'like tranny', 'look', 'look like', 'she', 'she look', 'tranny']
```

Named entity removing

```
wordnet_lemmatizer = WordNetLemmatizer()  
  
clean1 = d["tweet"].str.replace((r'@[\w]*'), '')  
clean2 = clean1.str.replace(r'RT', '')  
cleanTweets = clean2.str.replace(r'^a-zA-Z +^', '')  
cleanTweets[0]
```

```
' As a woman you shouldnt complain about cleaning up your house amp as a man you should always take the trash out'
```

Stopwords removing

```
lower_case = [[x.casefold() for x in sublst] for sublst in tokens]  
cleanTweets = [[wordsub for wordsub in word if wordsub not in stop_words] for word in lower_case]  
print(clean[2822])  
print(cleanTweets[2])
```

```
I wouldve won this tourney but then faggot ass Roy  
['wouldve', 'tourney', 'faggot', 'ass', 'roy']
```



Transformación de los datos

CountVectorizer


```
from sklearn.feature_extraction.text import CountVectorizer
count_vectorizer = CountVectorizer(ngram_range=(1,2))

countp = count_vectorizer.fit_transform(rand[3:5])
print(rand[3:5])
print(countp.toarray())
```

```
['I guess this is the night bitches die Stewie is that nigga', 'Sometimes I feel like being really nice to everyone then Im
just like wait youre all bitches anyway soooo']
[[0 0 0 0 0 0 1 0 1 1 1 0 0 0 0 1 1 0 0 2 1 1 0 0 0 0 0 0 0 1 1 1 0 0 0 0
 0 1 1 1 1 1 1 0 0 1 1 0 0 0 0 0 0]
[1 1 1 1 1 1 1 1 0 0 0 1 1 1 1 0 0 1 1 0 0 0 1 1 2 1 1 1 1 0 0 0 1 1 1 1
1 0 0 0 0 0 0 1 1 0 0 1 1 1 1 1 1]]
```




Exploración del dataset







Exploración del dataset





Exploración del dataset





Clasificación

1

Cross validation



MultinomialNB

alpha= 0.5

cv= KFold(10, shuffle=True)

85.83% accuracy (balanceado)

74.25% accuracy (desbalanceado)



RandomForestClassifier

cv= StratifiedKFold(3, shuffle=True)

92.21 % accuracy (balanceado)

83.95% accuracy (desbalanceado)



2

One-Vs-the-Rest



SVC

probability=True
kernel= 'linear'

86.28% accuracy (balanceado)

88.84% accuracy (desbalanceado)



MultinomialNB

alpha= 0.5

83.11 % accuracy (balanceado)

79.01% accuracy (desbalanceado)



3

Red neuronal



Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 10)	961150
dense_1 (Dense)	(None, 1)	11

Total params: 961,161

Trainable params: 961,161

Non-trainable params: 0






Librerías

matplotlib





Gracias!

¿Preguntas?

Repositorio Github

- <https://github.com/vgoyenechec/Hatespeech-en-twitter>